

AUDIRA: PROMPT CHAIN & LLM LOGIC FLOW

Version: 1.0

Purpose: Technical explanation of how Audira uses LLMs in a multi-step pipeline to transform business input into fully configured AI agents.

Overview

Audira uses a structured prompt chain consisting of **three primary stages** to convert onboarding answers into personalized dynamic questions, tag coverage, tone profiles, and simulated agent behavior.

Each stage uses controlled LLM instructions, tag context, and business metadata to keep the generation scoped and business-aligned.

STAGE 1: BUSINESS PROFILE SUMMARIZATION

Purpose:

Condense responses from the 10 fixed onboarding questions into a structured profile with three fields:

- **Business Objective**
- **Audience & Channel Context**
- **Industry Classification**

Prompt Template:

Summarize the client's business based on the following answers:

1. Core goal
2. Audience and tone
3. Channels to deploy on
4. Business outcomes expected
5. Compliance or training materials

Output as:

- Business Objective:
- Audience & Channel:
- Industry Context:

Output Example:

- Business Objective: Qualify property leads and schedule site visits.
 - Audience & Channel: Home buyers (30–55), using WhatsApp and website chat.
 - Industry Context: Real estate with moderate complexity and tone flexibility.
-

② STAGE 2: TAG ASSIGNMENT ENGINE

Purpose:

Apply the 10 discovery tags based on inferred context from the summarized profile.

Prompt Template:

Based on this profile, assign all relevant configuration tags from the list:

```
#intent_clarity, #channel_behavior, #tone_variation, #fallback_rules, #integration_scope,  
#training_data_source, #compliance_constraints, #persona_control,  
#recommendation_logic, #booking_handling
```

Explain why each tag is triggered in 1 sentence.

Output Example:

```
{  
  "#intent_clarity": "Client wants to qualify leads and book viewings.",  
  "#channel_behavior": "They use WhatsApp and website chat.",  
  "#booking_handling": "Site visit scheduling is a primary outcome."  
}
```

③ STAGE 3: DYNAMIC QUESTION GENERATOR

◆ Purpose:

Generate 10 tailored questions to clarify or finalize config based on missing or ambiguous tag logic.

✎ Prompt Template:

Using the tags and business profile, generate 10 dynamic discovery questions to:

1. Clarify missing behaviors
2. Finalize tone or fallback logic
3. Prepare integration planning
4. Reduce ambiguity before simulation

Each question should target a tag and include a short rationale.

📄 Output Example:

```
[  
 {  
   "question": "Should the AI suggest listings based on location or user preferences?",  
   "tag": "#recommendation_logic",  
   "reason": "Clarifies recommendation rules for real estate context."  
 },  
 {  
   "question": "What happens if a user asks to book on a national holiday?",  
   "tag": "#booking_handling",  
   "reason": "Helps finalize booking edge-case handling."  
 }]  


---


```

STAGE 4: DEEP-DIVE CLARIFIER

Triggered only if validation shows missing tags or risk/conflict (compliance, tone, fallback).

Uses a narrowed prompt:

Based on gaps in discovery tags, ask targeted questions to finalize configuration. Max 5 follow-ups.

Chain Control & Guardrails

- Temperature: 0.3–0.5 (controlled creativity)
 - Max tokens: 300–600 per stage
 - Custom validators check output for:
 - Tag presence
 - Question uniqueness
 - Alignment with profile context
-

Compatible Models

- gpt-4o, gpt-4, claude-3-opus, mistral-7b-instruct, llama3-70b-chat
 - Future support for in-house LLM via HuggingFace interface
-

Next Document: Pre-Launch Validator Spec