

## FSD\_09 – Admin Review Dashboard + Feedback Router

---

### Purpose:

This module provides **Audira admins and partner consultants** with a secure, centralized interface for:

- Reviewing agent readiness
- Overriding failed validator outputs
- Managing flagged memory or conflict cases
- Routing feedback and user corrections into the retraining pipeline

It acts as the **human-in-the-loop oversight layer** and feedback collection system for quality control and continuous learning.

---

### FSD\_09 – Section Breakdown

Section	Description
<b>1. Scope</b>	What the admin dashboard and feedback system are responsible for
<b>2. Input Sources</b>	What data feeds into the dashboard and router
<b>3. Dashboard Views &amp; Controls</b>	What admins can see, do, override, or flag
<b>4. Feedback Collection Points</b>	Where feedback is gathered (users, partners, QA)
<b>5. Feedback Processing Logic</b>	How input is interpreted, stored, and escalated
<b>6. Escalation &amp; Triage Logic</b>	How issues are resolved or routed internally
<b>7. Retraining Hooks</b>	How confirmed corrections are fed into agent learning or retraining
<b>8. Future Enhancements</b>	Audit trail visibility, auto-labeling, QA scoring, and agent benchmarking

## ◆ Section 1: Scope

---

### Purpose:

This module provides the **control panel** for internal reviewers, consultants, or authorized partners to:

- Inspect the full onboarding results of any agent
- View readiness status and validator output
- Manually override failed launch checks with justification
- Resolve memory conflicts or contradictions
- Route flagged issues to the correct team (content, QA, retraining)
- Collect structured feedback from users, testers, or field partners

It is **not a passive viewer** — it is an active governance interface in the Audira platform lifecycle.

---

### Responsibilities:

Function	Description
<b>Agent Readiness Review</b>	View detailed validator reports (from FSD_05), tag coverage, segment usage
<b>Override Gatekeeper</b>	Let authorized users force agent launch with justification
<b>Memory Conflict Manager</b>	Inspect, resolve, or flag contradictory agent facts (from FSD_07)
<b>Feedback Router</b>	Collect and triage user or QA feedback during or after agent use
<b>Retraining Trigger Engine</b>	Flag agents or blocks for AI model fine-tuning based on structured signals
<b>Audit &amp; Compliance Logging</b>	Maintain full trail of changes, overrides, and decision history

---

## Not in Scope:

Excluded Feature	Reason
<b>Editing onboarding answers</b>	That occurs in the onboarding interface
<b>Modifying uploaded documents</b>	Those are read-only once parsed in FSD_01
<b>Direct agent response generation</b>	That is handled live in FSD_08
<b>LLM prompting or inference logic</b>	Controlled by FSD_06 and FSD_08

## Supported Reviewer Roles:

Role	Capabilities
<b>Internal Admin</b>	Full access to override, resolve, revalidate, and mark agents for retraining
<b>Partner Consultant</b>	Read + resolve permissions only for assigned clients
<b>QA Team Member</b>	Can simulate agent behavior, submit scoring feedback
<b>Feedback Analyst</b>	Can review user-captured issues, but not resolve conflicts

## References:

-  *AUDIRA PRE-LAUNCH VALIDATOR SPEC* – defines readiness report and override logic
-  *AUDIRA AGENT BLUEPRINT TEMPLATE* – identifies agent role, owner, partner metadata
-  *AUDIRA FILE & DATA UPLOAD SCHEMA* – allows source tracing for tag conflicts
-  *AUDIRA PROMPT CHAIN & LLM LOGIC FLOW* – provides context for simulation/test mode
-  *AUDIRA PRODUCT BLUEPRINT* – defines governance, QA, and escalation process roles

## ◆ Section 2: Input Sources

This section defines all the data streams that feed into the Admin Review Dashboard and Feedback Router. These inputs are pulled from earlier FSD modules and real-time interaction logs.

---

### A. Readiness Validator Output

**Source:** FSD\_05

Provides overall launch readiness, score, failed areas, and auto-suggested fixes.

```
{  
  "agent_id": "AGENT_00231",  
  "readiness_score": 71,  
  "status": "FAIL",  
  "failed_areas": [  
    "MISSING_CRITICAL_TAG:pricing_strategy",  
    "DOCUMENT_GAPS:recency"  
  ],  
  "suggested_fixes": [  
    "Upload a recent overview doc",  
    "Answer onboarding question Q3"  
  ]  
}
```

---

### B. Memory Conflict Records

**Source:** FSD\_07

Captures unresolved contradictions between user answers, document content, or segment analysis.

```
{
```

```
"tag_id": "payout_policy",
"conflict_detected": true,
"status": "pending_resolution",
"values": ["automated via Stripe", "manual payout by finance"]
}
```

---

### C. Agent Metadata & Intent

**Source:** Agent Onboarding Framework

Includes industry, launch type, partner owner, and module capabilities.

```
{
  "agent_id": "AGENT_00231",
  "industry": "SaaS",
  "configured_by": "partner_consultant_127",
  "capabilities": ["Q&A", "Sales"]
}
```

---

### D. Feedback Logs

**Sources:**

- Agent chat interface (user rating buttons, “Was this helpful?”)
- QA test results
- Admin simulations (FSD\_10 preview hooks)

```
{
  "feedback_id": "fbk_90211",
  "source": "user_live_session",
  "linked_response": "Agent said: Vendors are paid manually.",
  "feedback_type": "incorrect_info",
```

```
"submitted_by": "user_342",
"timestamp": "2025-06-11T03:33:00Z"
}
```

---

## E. Prompt & Response Trace Logs

**Source:** FSD\_08

Used for simulation replay and traceable decision flow.

```
{
  "turn_id": "turn_8181",
  "model_used": "gpt-4",
  "input_prompt": "...",
  "response": "You offer three pricing tiers...",
  "tags_covered": ["pricing_model"],
  "memory_used": true
}
```

---

## Summary Input Table

Input	Used For
<b>readiness_report</b>	Validator decision review and override logic
<b>conflict_log</b>	Memory inspection and contradiction resolution
<b>agent_metadata</b>	Display filters and partner-specific views
<b>feedback_items</b>	Triage for QA, retraining, or correction
<b>response_traces</b>	Replay and response behavior review

---

## References:

-  *AUDIRA PRE-LAUNCH VALIDATOR SPEC* – defines readiness scoring fields and logic
  -  *AUDIRA FILE & DATA UPLOAD SCHEMA* – provides document + segment anchors
  -  *AUDIRA PROMPT CHAIN & LLM LOGIC FLOW* – used to generate replay paths
  -  *AUDIRA AGENT ONBOARDING FRAMEWORK* – sets up reviewer access rights
- 

## Section 3: Dashboard Views & Controls

This section defines the **UI-level experience** for authorized reviewers, including what data is visible, what controls are available, and how each agent can be inspected or overridden.

---

### Core Views in the Admin Dashboard

View	Description
<b>Agent Overview Panel</b>	Summary of agent status, readiness, creator, last update
<b>Readiness Report Viewer</b>	Full breakdown of FSD_05 validator output
<b>Conflict Resolver Panel</b>	View, compare, and resolve memory conflicts (FSD_07)
<b>Prompt + Response Simulation</b>	Run LLM previews using current prompt/memory
<b>Feedback Inbox</b>	Sortable, filterable feedback queue from users/QA
<b>Override History Log</b>	List of all past admin interventions per agent

---

### Key Controls & Actions

Control	Description	Role Permissions

<b>Override Readiness Failure</b>	Force agent launch with justification	Admin only
<b>Resolve Conflict</b>	Select correct value or mark for follow-up	Admin + Partner Consultant
<b>Re-run Validator</b>	Recalculate FSD_05 after update	Admin only
<b>Simulate Agent Response</b>	Preview how the agent would reply	Admin + QA
<b>Route Feedback to Retraining</b>	Mark feedback for model update	Admin only
<b>Mark Agent as “Trusted”</b>	Locks prompt/memory from AI overwrite	Admin only
<b>Flag Segment or Tag</b>	Mark source content as unreliable or incorrect	Any reviewer

## Dashboard Filters

Filter Option	Description
Agent status (e.g., “failed”, “ready”, “launched”)	
Industry (e.g., retail, SaaS, health)	
Tag conflict count	
Last modified (date range)	
Partner owner	
“Has unresolved feedback” flag	

## Smart Routing Widgets (Optional AI Boost)

Widget	Function
<b>Suggest Override Reasons</b>	LLM recommends justification text based on validator gaps

<b>Auto-Summarize Feedback Trends</b>	Group feedback by type, frequency, tag impact
<b>Conflict Resolution Draft Generator</b>	LLM proposes merged or clarified tag value based on evidence

---

### Example View: Agent Readiness Panel

Field	Value
<b>Agent ID</b>	AGENT_00231
<b>Created By</b>	partner_consultant_127
<b>Readiness Score</b>	71 (Fail)
<b>Failed Areas</b>	MISSING_TAG:pricing_model, DOCUMENT_GAPS:overview
<b>Conflict Count</b>	2 unresolved
<b>Last Simulated Response</b>	“Support is 24/7 via Zendesk...”
<b>Override Button</b>	[Available]
<b>Conflict Panel</b>	[Resolve Now]
<b>Feedback Linked</b>	3 items

---

### References:

-  *AUDIRA AGENT BLUEPRINT TEMPLATE* – configures reviewer roles and view permissions
  -  *AUDIRA PRE-LAUNCH VALIDATOR SPEC* – feeds report view and override logic
  -  *AUDIRA PROMPT CHAIN & LLM LOGIC FLOW* – supports simulation and response injection
  -  *AUDIRA PRODUCT BLUEPRINT* – defines governance workflow and intervention policies
-

## ◆ Section 4: Feedback Collection Points

This section outlines where, when, and how structured feedback is collected from users, QA testers, and admins throughout the agent lifecycle.

---

### ❖ Key Feedback Collection Touchpoints

Source	Trigger Point	Type of Feedback
<b>End User (Live Agent UI)</b>	After each response or session	“Was this helpful?” thumbs up/down, comment
<b>QA Simulations</b>	During scripted test runs	Scorecards, tagged errors, simulated corrections
<b>Admin Review Actions</b>	When resolving conflicts or overriding logic	Comments and audit flags
<b>Partner Consultant Reviews</b>	Onboarding validation steps	Suggested changes, confidence notes
<b>Post-Deployment Audits</b>	Periodic quality evaluations	Graded response samples and segment quality scores

---

### 📁 Feedback Categories

Category	Example
<b>Incorrect Information</b>	“Agent said we offer 24/7 support, but we don’t.”
<b>Outdated Segment</b>	“Doc mentions old pricing. New plans launched in 2025.”
<b>Missing Answer</b>	“Agent didn’t know our return policy.”
<b>Inappropriate Tone</b>	“Too casual — should be professional in finance use case.”
<b>Hallucination Risk</b>	“Agent invented info not in memory or prompt.”
<b>Conflicting Memory</b>	“Two different answers about vendor onboarding.”

---

### 📝 Feedback Form Fields

Each piece of submitted feedback includes:

```
{  
  "feedback_id": "fbk_72819",  
  "agent_id": "AGENT_0123",  
  "source": "user",  
  "feedback_type": "incorrect_info",  
  "linked_tag": "payout_policy",  
  "linked_response": "We pay vendors every Friday via ACH.",  
  "comment": "Not true. We pay on the 15th and 30th via Stripe.",  
  "submitted_by": "user_342",  
  "timestamp": "2025-06-11T03:49:00Z"  
}
```

---

## Submission Channels

Channel	Form Type	Real-time Sync
In-agent chat interface	Pop-up form	<input checked="" type="checkbox"/>
Admin dashboard override	Text + tag selector	<input checked="" type="checkbox"/>
QA console	Structured test feedback	<input checked="" type="checkbox"/>
Partner portal	Free-text + optional prompt view	<input checked="" type="checkbox"/>

---

## References:

-  *AUDIRA AGENT ONBOARDING FRAMEWORK* – identifies who can submit feedback and when
-  *AUDIRA PROMPT CHAIN & LLM LOGIC FLOW* – used to link feedback to specific prompt + memory

-  **AUDIRA AGENT SIMULATION TEST KIT** – defines QA flow and feedback capture from scripted tests
  -  **AUDIRA FILE & DATA UPLOAD SCHEMA** – maps flagged segments or responses back to document source
- 

## ◆ **Section 5: Feedback Processing Logic**

This section defines **how feedback is evaluated, categorized, prioritized, and routed** inside the system — enabling both manual triage and automated actions when appropriate.

---

### **Feedback Intake Pipeline**

#### 1. **Feedback Submission**

→ From live user, QA team, admin, or partner.

#### 2. **Auto-Categorization**

→ NLP-based classification + tag detection (e.g., classify as incorrect\_info, link to pricing\_model).

#### 3. **Confidence Scoring**

→ Based on:

- Feedback source role (e.g., user vs. admin)
- Text similarity to memory or prompt
- Frequency of similar reports for the same tag/segment

#### 4. **Routing Decision**

→ Sent to:

- Manual triage queue (for critical/conflicting items)
- Retraining candidate list (if confirmed)
- Admin reviewer (if tagged as override-worthy)

#### 5. **Linkage Storage**

→ Logs full context:

- Agent memory at time of issue

- Prompt block used
  - Segment or tag involved
  - Who submitted the feedback
- 

## Prioritization Rules

Priority Level	Trigger Criteria
High	Feedback affects core tag, has multiple reports, or comes from QA/admin
Medium	Single-source feedback with confirmed link to active prompt or memory
Low	Vague, low-confidence feedback from end users or on optional tags

---

## Example Processed Feedback Object

```
{  
  "feedback_id": "fbk_00213",  
  "tag_id": "vendor_model",  
  "feedback_type": "conflicting_memory",  
  "priority": "high",  
  "routing_path": ["admin_review", "retraining_queue"],  
  "linked_prompt": "AGENT_101_prompt_v4",  
  "linked_memory": "vendor_model_v2",  
  "status": "open",  
  "auto_classification_score": 0.93  
}
```

---

## System Actions Based on Feedback

Action	Trigger
<b>Flag Memory for Review</b>	If feedback contradicts an active memory block
<b>Mark Segment as Outdated</b>	If flagged for being old, incorrect, or misleading
<b>Soft-Disable Prompt Block</b>	If associated prompt fragment is triggering confusion
<b>Add to Retraining Set</b>	If feedback passes trust threshold or is admin-approved
<b>Notify Agent Owner</b>	For critical issues in live agents
<b>Trigger Follow-up Question</b>	Sends a new clarification query to the user (via FSD_04)

---

### References:

-  *AUDIRA PRE-LAUNCH VALIDATOR SPEC* – defines core vs. optional tag priority
-  *AUDIRA PROMPT CHAIN & LLM LOGIC FLOW* – maps feedback to generated prompt structures
-  *AUDIRA AGENT SIMULATION TEST KIT* – used to compare QA feedback against expected response paths
-  *AUDIRA AGENT BLUEPRINT TEMPLATE* – identifies roles with authority to approve feedback-triggered retraining

---

### Section 6: Escalation & Triage Logic

This section defines **how flagged issues, feedback, and unresolved memory conflicts are escalated** to the appropriate team or reviewer — ensuring every critical item gets resolved or documented.

---

### Escalation Triggers

Escalation Trigger	Routed To
 <b>Repeated user flags on same tag or response</b>	QA team + admin reviewer
 <b>Memory conflict unresolved after override</b>	Memory team or AI engineer

Agent answering with outdated or hallucinated info	Retraining lead + owner
Failed validator override with risk tags	Executive reviewer with justification log
Feedback on compliance, legal, or health tags	Legal/partner compliance team
Ambiguous multi-source segment disagreement	Sent to document integrity triage queue

---

### Issue Routing Logic

Each flagged item is assigned:

- issue\_type: e.g., conflict, hallucination, outdated info
  - severity\_level: auto-detected or reviewer-assigned
  - assigned\_to: team or individual based on tag/category ownership
  - status: open / under review / resolved / retraining scheduled
  - escalation\_trace: full record of who handled what and when
- 

### Escalation Panel in Dashboard

Field	Description
<b>Issue Type</b>	Memory conflict, bad response, prompt mismatch
<b>Affected Tags</b>	One or more discovery tags linked
<b>Affected Agent</b>	ID, owner, launch date
<b>Last Action</b>	E.g., “Manually resolved by QA” or “Pending retraining”
<b>Days Open</b>	SLA tracking for unresolved issues
<b>Status</b>	Open, Assigned, Resolved, Archived

---

## Auto-Loop to Retraining or Validator

If resolution involves factual change:

- Agent memory is patched (FSD\_07)
- Prompt chain may be recompiled (FSD\_06)
- Readiness validator is optionally re-run (FSD\_05)

For major issues (e.g., hallucination from prompt logic), the escalation may trigger a **prompt chain logic audit**.

---

## Audit Trail Requirements

Each escalated case must maintain:

- Feedback text or source trace
  - Action history (resolves, rejections, reassignment)
  - If override used: justification + reviewer ID
  - Snapshot of memory and prompt at time of issue
- 

## References:

-  *AUDIRA PRODUCT BLUEPRINT* – defines escalation thresholds and reviewer responsibilities
  -  *AUDIRA AGENT BLUEPRINT TEMPLATE* – maps agents to reviewers, QA teams, and risk level
  -  *AUDIRA FILE & DATA UPLOAD SCHEMA* – allows document-linked triage and escalation
  -  *AUDIRA PROMPT CHAIN & LLM LOGIC FLOW* – provides structure for audit snapshot reconstruction
-

## ◆ Section 7: Retraining Hooks

This section defines **how validated feedback, memory updates, and escalated cases are used to trigger targeted retraining** of LLM agents — ensuring continuous improvement and alignment with each SMB's evolving profile.

---

### ⌚ What Can Trigger Retraining?

Trigger Type	Description
<input checked="" type="checkbox"/> <b>Admin-confirmed feedback</b>	Reviewed and approved correction or clarification
<input checked="" type="checkbox"/> <b>Overridden memory conflict</b>	Once a contradiction is resolved and marked as “final”
<input checked="" type="checkbox"/> <b>High-volume user feedback</b>	Same tag or response flagged by multiple users or sessions
<input checked="" type="checkbox"/> <b>QA flagged test failure</b>	Tagged with retraining suggestion during simulation
<input checked="" type="checkbox"/> <b>Prompt-to-response mismatch</b>	Validated case of prompt block resulting in hallucination
<input checked="" type="checkbox"/> <b>Agent delta after re-onboarding</b>	When documents or answers change post-launch

---

### 📁 Retraining Data Format

All training candidates are compiled as `TrainingSample` objects:

```
{  
    "agent_id": "AGENT_1021",  
    "issue_type": "conflicting_memory",  
    "corrected_tag": "payout_policy",  
    "correct_value": "Automated via Stripe every 14 days",  
    "source": "admin_resolution",  
    "linked_prompt": "prompt_v5",
```

```
"llm_response": "Manual payout confirmed",
"training_action": "negative_sample",
"timestamp": "2025-06-11T04:00:00Z"
}
```

---

## Training Actions

Action Type	Usage
<b>Positive Sample</b>	“This is what the agent should say in this context.”
<b>Negative Sample</b>	“This is what the agent said incorrectly — avoid this.”
<b>Replacement Sample</b>	“Replace this behavior with updated memory or tone.”

---

## Retraining Hooks Queue

Queue	Triggered By	Processing Mode
<b>training/live_corrections</b>	Resolved feedback from users/admin	Batched nightly
<b>training/conflict_resolution</b>	Memory contradictions resolved	Continuous
<b>training/qa_failures</b>	Failed simulations	Weekly model update round
<b>training/hallucination_cases</b>	Prompt logic overrides	Priority retrain window
<b>training/new_tags</b>	New tag types or values not yet modeled	Tag dictionary + prompt template updates

---

## Evaluation Before Push

Before integrating into model fine-tuning:

- Each sample is evaluated for clarity and source traceability
- Language and tone are normalized

- Examples are tested in **agent simulation preview mode** (FSD\_10)
  - Approved samples are added to the training corpus per agent class (e.g., “Retail SME agent v3.2”)
- 

## References:

-  *AUDIRA AGENT SIMULATION TEST KIT* – used for retraining sample QA before LLM fine-tuning
  -  *AUDIRA PROMPT CHAIN & LLM LOGIC FLOW* – needed for prompt-response alignment during model update
  -  *AUDIRA AGENT ONBOARDING FRAMEWORK* – defines class-level agent groupings for training sets
  -  *AUDIRA DISCOVERY TAGS DICTIONARY* – expanded when training introduces new or rare tag patterns
- 

## Section 8: Future Enhancements

This section outlines strategic improvements and tooling extensions planned to evolve the dashboard and feedback loop into a full **QA governance hub** for Audira and its partners.

---

## Upcoming Features & Enhancements

Feature	Description	Benefit
<b>Real-Time Feedback Streaming</b>	View live thumbs up/down and user comments as they happen	Enables instant QA visibility
<b>Audit Trail Explorer</b>	View full lifecycle of memory/prompt changes, overrides, and response deltas	Improves traceability and compliance
<b>Feedback Heatmaps</b>	Visualize which tags or topics are triggering the most flags or corrections	Identifies knowledge gaps at scale

<b>Auto-Labeling with LLMs</b>	Use AI to pre-label feedback and classify root causes	Saves human triage time
<b>Reviewer Scoring Dashboard</b>	Tracks reviewer activity, resolution times, accuracy of overrides	Enables internal performance metrics
<b>Agent Benchmark Simulator</b>	Periodic test harness for launched agents, scored against a standard QA set	Enables automated regression testing
<b>Time-Sensitive Memory Reminders</b>	Prompts admins to review memory blocks flagged as stale or policy-expired	Improves memory freshness
<b>Custom QA Templates by Industry</b>	Tailor review forms and escalation rules per vertical (e.g., fintech vs. health)	Enhances reviewer relevance
<b>Feedback Reuse Suggestions</b>	Surface similar past issues and resolution paths when a new feedback case matches prior ones	Boosts reviewer efficiency

## Compatible OSS Tools

Tool	Role
<b>TruLens / DeepEval</b>	Feedback scoring and fine-tuning evaluation
<b>LangGraph</b>	Issue triage routing and override workflows
<b>PromptLayer / PromptTools</b>	Prompt-response audit versioning
<b>Label Studio</b>	Optional front-end for feedback classification
<b>Weights &amp; Biases</b>	Logging model changes triggered by dashboard feedback
<b>Unstructured.io</b>	Segment mapping for flagged source fragments

## Linked Modules:

-  *AUDIRA PRODUCT BLUEPRINT* – sets roadmap for feedback-to-training evolution
-  *AUDIRA PROMPT CHAIN & LLM LOGIC FLOW* – enables response-to-feedback mapping

-  *AUDIRA PRE-LAUNCH VALIDATOR SPEC* – source of override, scoring, and readiness data
  -  *AUDIRA AGENT BLUEPRINT TEMPLATE* – defines reviewer roles, escalation privileges, and agent groupings
-